


Learn Torch in 60 mins

Kai Zhao James Cross Liang Huang
Dept. EECS, Oregon State University



Teach yourself variables, constants, arrays, strings, expressions, statements, functions,...



Teach yourself program flow, pointers, references, classes, objects, inheritance, polymorphism,



Do a lot of recreational programming. Have fun hacking but remember to learn from your mistakes.



Interact with other programmers. Work on programming projects together. Learn from them.



Teach yourself advanced theoretical physics and formulate a consistent theory of quantum gravity.




**Teach yourself biochemistry,
molecular biology, genetics,...**



Use knowledge of biology to make an age-reversing potion.



Use knowledge of physics to build flux capacitor and go back in time to day 21.



Replace younger self.



As far as I know, this is the easiest way to "Teach Yourself C++ in 21 Days".

Toolkits in the Thriving Deep Learning Community



Schedule of the Tutorial

Torch (Today)

Provides high-level abstractions
as well as low-level access



DyNet (Dec. 2nd)

Specialized for dynamically
changing networks

TensorFlow (Dec. 9th)

Industrial level toolkit
Support massive GPU clusters



In Today's Tutorial

- Basics of Lua Language
- Basic Tensor Operations in Torch
- Building a Simple Network
- Training the Simple Network
- Building a Recurrent Neural Network


In Today's Tutorial

- Basics of Lua Language
- Basic Tensor Operations in Torch
- Building a Simple Network
- Training the Simple Network
- Building a Recurrent Neural Network

Lua: Simple & Fast

- Designed as an embedded language. (e.g., in *World of Warcraft*)
 - Extremely simple grammar:
 - Atomic data types: float, string, boolean
 - Everything else is a hash table
 - First-class functions
 - Surprisingly fast as a script language:
 - LuaJIT is very well developed
 - Speed comparable to Java
 - Easy interaction with C/C++
 - Simple interface
 - Little overhead

Learn Lua for the Horde!



Lua: Data Types & Flow Control

- Variables, Flow Control, and Functions
 - Same as most imperative languages
 - Every undefined variable is by default *nil*
- Only one compound data type: *table*

```
-- Dict literals have string keys by default:
t = {key1 = 'value1', key2 = false}

-- String keys can use js-like dot notation:
print(t.key1)  -- Prints 'value1'.
t.newKey = {}  -- Adds a new key/value pair.
t.key2 = nil   -- Removes key2 from the table.

-- Literal notation for any (non-nil) value as key:
u = {'@!#' = 'qbert', [{]} = 1729, [6.28] = 'tau'}
print(u[6.28]) -- prints "tau"

-- Key matching is basically by value for numbers and strings,
but by identity
-- for tables.
a = u['@!#']  -- Now a = 'qbert'.
b = u[{]}     -- We might expect 1729, but it's nil
```

Lua: Data Types & Flow Control

- Variables, Flow Control, and Functions
 - Same as most imperative languages
 - Every undefined variable is by default *nil*
- Only one compound data type: *table*
 - List/Array: table w/ consecutive int. keys (index from 1)

```
-- List literals implicitly set up int keys:  
v = {'value1', 'value2', 1.21, 'gigawatts'}  
for i = 1, #v do -- #v is the size of v for lists.  
    print(v[i]) -- Indices start at 1 !! SO CRAZY!  
end  
-- A 'list' is not a real type. v is just a table with  
-- consecutive integer keys, treated as a list.
```

- Iterate through table

```
for key, val in pairs(u) do -- Table iteration.  
    print(key, val)  
end
```

Lua: OOP

- Class is just another table

```
Dog = {} -- 1.

function Dog:new() -- 2.
    local newObj = {sound = 'woof'} -- 3.
    self.__index = self -- 4.
    return setmetatable(newObj, self) -- 5.
end

function Dog:makeSound() -- 6.
    print('I say ' .. self.sound)
end

mrDog = Dog:new() -- 7.
mrDog:makeSound() -- 'I say woof' -- 8.
```

- Definition member function

function tablename:fn(...) ... **end**

equals to

function tablename.fn(self, ...) ... **end**

In Today's Tutorial

- Basics of Lua Language
- Basic Tensor Operations in Torch
- Building a Simple Network
- Training the Simple Network
- Building a Recurrent Neural Network

Torch: Tensors

- Basic operands: Tensors

- Types: ByteTensor, CharTensor, ShortTensor, IntTensor, LongTensor, FloatTensor, DoubleTensor

```
a = torch.Tensor(5, 3)  -- construct a 5x3 matrix uninitialized

a = torch.rand(5, 3)
print(a)
```

- Simple Operators

- E.g., multiplication

```
-- matrix-matrix multiplication: syntax 1
a*b

-- matrix-matrix multiplication: syntax 2
torch.mm(a, b)

-- matrix-matrix multiplication: syntax 3
c = torch.Tensor(5, 4)
c:mm(a, b)  -- store the result in c
```

Torch: Other Operations

- Constructors: `torch.ones()` `torch.zeros()`
- Element-wise Operators: `abs()`, `pow()`
- Column-wise Operators: `sum()`, `max()`
- Matrix-wise Operators: `trace()`, `norm()`

```
torch.cat(torch.ones(3), torch.zeros(2))  
1  
1  
1  
0  
0  
[torch.DoubleTensor of size 5]
```

```
torch.cat(torch.ones(3, 2), torch.zeros(2, 2), 1)  
1 1  
1 1  
1 1  
0 0  
0 0  
[torch.DoubleTensor of size 5]
```

In Today's Tutorial

- Basics of Lua Language
- Basic Tensor Operations in Torch
- **Building a Simple Network**
- Training the Simple Network
- Building a Recurrent Neural Network

Neural Networks

- Package 'nn'
 - Basic neural network modules
 - Construction methods
- Linear module as an example

```
require 'nn';  
lin = nn.Linear(5, 3)
```

- Just another table

```
lin  
nn.Linear(5 -> 3)  
{  
  gradBias : DoubleTensor - size: 3  
  weight : DoubleTensor - size: 3x5  
  _type : torch.DoubleTensor  
  output : DoubleTensor - empty  
  gradInput : DoubleTensor - empty  
  bias : DoubleTensor - size: 3  
  gradWeight : DoubleTensor - size: 3x5  
}
```

output = weight * X + bias

Neural Networks

- Forward/Backward already defined for modules

- Forward

```
y = lin.forward(x)
print(y)
```

- Backward

```
lin.backward(x, grad)
```

- Call `:zeroGradParameters()` before backward

- Now we can manually do gradient descent

```
lin.weight.add(0.1*lin.gradWeight)
lin.bias.add(0.1*lin.gradBias)
```

I believe in graduate student descent.

— David McAllester



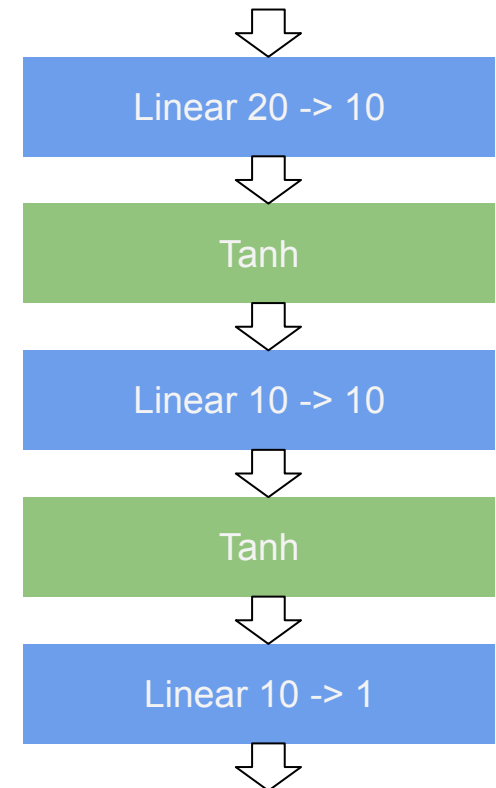
Neural Networks

- Composing more complicated networks

- Use package 'nn'

```
net = nn.Sequential();  
net.add(nn.Linear(20, 10));  
net.add(nn.Tanh());  
net.add(nn.Linear(10, 10));  
net.add(nn.Tanh());  
net.add(nn.Linear(10, 1));
```

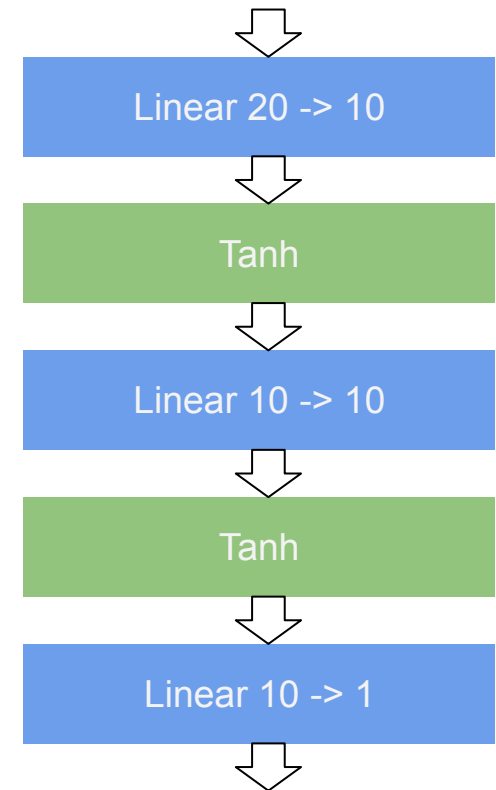
```
x = torch.rand(20)  
y1 = net.forward(x)  
print(y1)  
-0.2648  
[torch.DoubleTensor of size 1]
```



Neural Networks

- Composing more complicated networks
 - Use package 'nngraph'

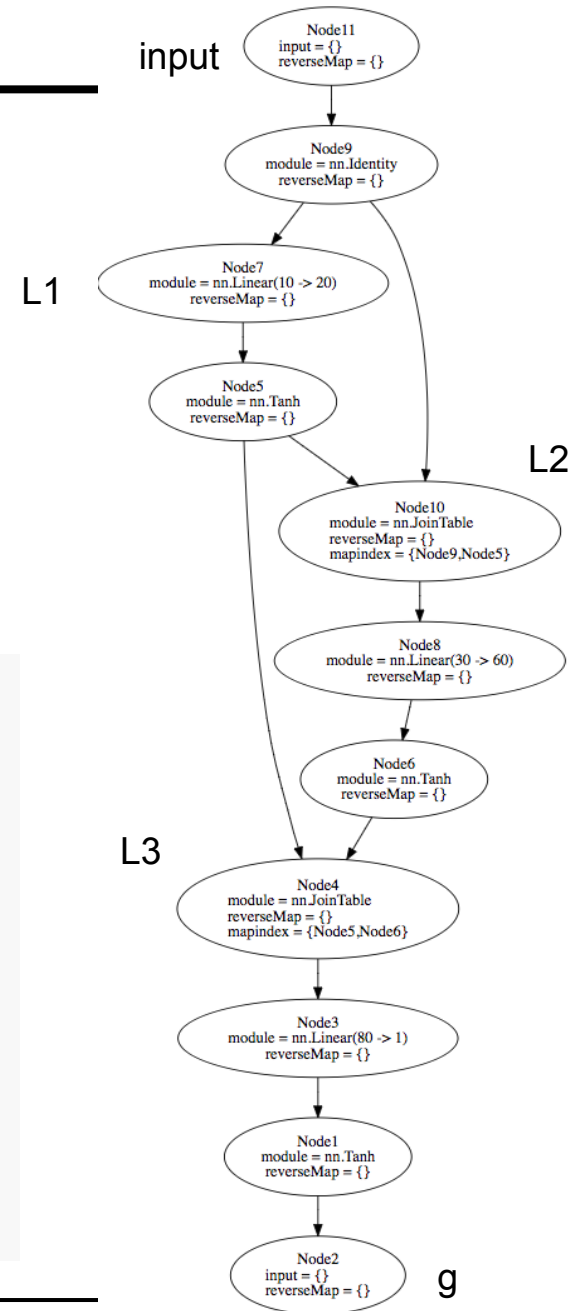
```
require 'nngraph';  
g1 = - nn.Linear(20, 10)  
  
g2 = g1  
  - nn.Tanh()  
  - nn.Linear(10, 10)  
  - nn.Tanh()  
  - nn.Linear(10, 1)  
gnet = nn.gModule({g1}, {g2})
```



Neural Networks

- Composing more complicated networks
 - 'nngraph' is easier to use than 'nn'

```
input = - nn.Identity()
L1 = input
    - nn.Linear(10, 20)
    - nn.Tanh()
L2 = {input, L1}
    - nn.JoinTable(1)
    - nn.Linear(30, 60)
    - nn.Tanh()
L3 = {L1, L2}
    - nn.JoinTable(1)
    - nn.Linear(80, 1)
    - nn.Tanh()
g = nn.gModule({input},{L3})
```



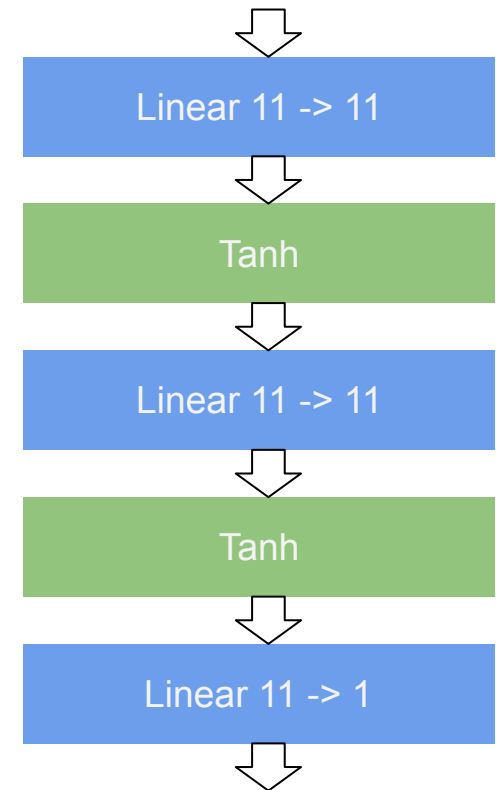
In Today's Tutorial

- Basics of Lua Language
- Basic Tensor Operations in Torch
- Building a Simple Network
- **Training the Simple Network**
- Building a Recurrent Neural Network

Training: Setup

- Dataset: UCI red wine quality
 - X: 11 dims real numbers in [0,1]
 - Y: real number in [0, 1]
- Loss: Mean Square Error

```
torch.manualSeed(1234)
-- build the network
g1 = - nn.Linear(11, 11)
g2 = g1
    - nn.Tanh()
    - nn.Linear(11, 11)
    - nn.Tanh()
    - nn.Linear(11, 1)
winenet = nn.gModule({g1}, {g2})
-- mean square error
loss = nn.MSECriterion()
```



Training: General Framework

```
W, gradW = winenet:getParameters()  
optimState = {}
```

get the location of the weights
and the grad weights

```
for epoch = 1, n_epochs do
```

```
  local total_loss = 0
```

```
  for i=1, n_examples do
```

```
    x = xx[i]
```

```
    y = torch.Tensor({yy[i]})
```

```
    winenet:zeroGradParameters()
```

clean accumulated grad weights

```
    function feval()
```

```
      local predicted = winenet:forward(x)
```

forward through network

```
      local L = loss:forward(predicted, y)
```

calculate loss

```
      total_loss = total_loss + L
```

```
      local dL_dy = loss:backward(predicted, y)
```

get grad from loss

```
      winenet:backward(x, dL_dy) -- computes and updates gradW
```

backward through network

```
      return L, gradW
```

```
    end
```

```
    optim.sgd(feval, W, optimState)
```

update weights using
accumulated grad weights

```
  end
```

```
  print('at epoch', epoch, 'avg loss', total_loss/n_examples)
```


```
end
```

In Today's Tutorial

- Basics of Lua Language
- Basic Tensor Operations in Torch
- Building a Simple Network
- Training the Simple Network
- Building a Recurrent Neural Network

A Simple Recurrent Network


- The input of the network is a sequence of vectors:
 $x_t, t = 1 \dots n$
- At each step we recurrently calculate function:
$$h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$$
- Parameters W_x, W_h , and b are shared in different steps
- Size of the network varies with the input length



A Simple Recurrent Network

- Building block: $h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$

```
ht1 = - nn.Identity()  
xt = - nn.Identity()  
ht = {ht1, xt}  
    - nn.JoinTable(1)  
    - nn.Linear(20, 10)  
    - nn.Tanh()  
stepfunction = nn.gModule({ht1, xt}, {ht})
```



A Simple Recurrent Network

- Clone step function many times
- All clones **share** the same parameter at every time step
 - point to the **same** memory

```
W, gradW = stepfunction:getParameters()
```

```
clones = {}
```

```
N = 100
```

```
for i = 1, N do
```


```
    clones[i] = stepfunction:clone()
```

```
    share_params(clones[i], stepfunction)
```

```
end
```

**getParameters() reallocate memory
do it before cloning**

skipped for space reason, refer to ipython notebook




A Simple Recurrent Network

- Forward through time:

```
n = 7
h0 = torch.rand(10)
x = torch.rand(n, 10)
h = torch.zeros(n+1, 10)

W:uniform(-1, 1)


-- forward
h[1] = h0
for i = 1, n do
    h[i+1] = clones[i]:forward{h[i], x[i]}
end
```



A Simple Recurrent Network


- Backward through time:

```
stepfunction:zeroGradParameters()  
  
grad_h = torch.zeros(h:size())  
grad_h[n+1] = torch.rand(10)  
for i = n, 1, -1 do  
    local grads = clones[i]:backward({h[i], x[i]}, grad_h[i+1])  
    grad_h[i], grad_xi = unpack(grads)  
end
```



A Simple Recurrent Network

- Simple RNNs sometimes are not very effective
 - Vanishing gradient problem
- LSTM/GRU become more popular
 - Several packages contain these modules
- We do not cover them in this tutorial
 - But we demo an example (from Andrej Karpathy)
 - Text Generation trained on Shakespeare's works



FIN

thanks!

Next Session: Dec 2, 2016

Deep Learning with DyNet

References

- *Learn X in Y minutes, X=Lua*
<https://learnxinyminutes.com/docs/lua/>
- *Deep Learning with Torch* in CVPR 2015
<https://github.com/soumith/cvpr2015/blob/master/Deep%20Learning%20with%20Torch.ipynb>
- *Torch Video Tutorials*
<https://github.com/Atcold/torch-Video-Tutorials>
- *The Unreasonable Effectiveness of Recurrent Neural Networks*
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>